

Problem Statement

(40 points) Image Segmentation using mean shift.

Input image: 3/data/baboonColor.png.

Take this 512 512 pixel image, smooth it using Gaussian convolution with standard deviation 1 pixel width, and subsample the smoothed image by a factor of 2 in each spatial dimension to produce a 256 256 image. Use this smaller-sized image for the following experiment. If this image is still too large for your computers memory, then you may resize further.

Implement the algorithm for mean-shift image segmentation using both color (RGB) and spatial coordinate (XY) features. Tune parameters suitably to get a segmented image with at least 5 segments and no more than 50 segments. To improve code efficiency, you may use Matlab functions like `knnsearch()`, `bsxfun()`, etc. For this image, about 20 iterations should be sufficient for reaching close to convergence. You may select a random subset of nearest neighbors, in feature space, for the mean-shift updates to reduce running time. Each iteration can run in about 10-20 seconds on a typical personal computer.

1. Write a function `myMeanShiftSegmentation.m` to implement this.
2. Display the (i) original image along with (ii) the segmented image that shows color-coded pixels (and, thus, segments) using the color component of the converged feature vectors.
3. Report the following parameter values: Gaussian kernel bandwidth for the color feature, Gaussian kernel bandwidth for the spatial feature, number of iterations.

Code

Mean Shift Segmentation

```
1 function [processedImage, outputImage] = myMeanShiftSegmentation
   (inputImage)
2
3     W = 2;
4     sigma1 = 0.66;
5     numIterations = 20;
6     spaceSigma = 20;
7     intensitySigma = 20;
8
9     sigmas = -2 .* ([spaceSigma, spaceSigma, intensitySigma,
10                     intensitySigma, intensitySigma] .^ 2);
11     sigmaProds = (2*pi)^2.5 * spaceSigma * spaceSigma *
12                 intensitySigma * intensitySigma * intensitySigma;
```

```

13
14     h = fspecial('gaussian', [W, W], sigma1);
15
16     for i = 1:3
17         processedImage(:, :, i) = myShrinkImageByFactorD
            (imfilter(myLinearContrastStretching(
                inputImage(:, :, i)), h), 2);
18     end
19
20     numPixels = size(processedImage, 1) * size(
        processedImage, 2);
21     imageRepresentation = zeros(numPixels, 5);
22     newImageRepresentation = zeros(size(imageRepresentation)
        );
23
24     % initialize imageRepresentation
25     count = 0;
26     for x = 1:size(processedImage, 1)
27         for y = 1:size(processedImage, 2)
28             count = count + 1;
29             imageRepresentation(count, :) = [x/255,
                y/255, processedImage(x, y, 1),
                processedImage(x, y, 2),
                processedImage(x, y, 3)];
30         end
31     end
32
33     for i = 1:numIterations
34         fprintf('Iter: %d\n', i);
35         [IDX, D] = knnsearch(imageRepresentation,
            imageRepresentation, 'K', k, 'IncludeTies', ~
            true); %handle cell situation later
36
37         for x = 1:size(processedImage, 1)
38             for y = 1:size(processedImage, 2)
39                 rowInRepresentationMatrix = (x
                    -1) * size(processedImage, 2)
                    + y;
40
41                 points = IDX(
                    rowInRepresentationMatrix, 2:
                    end); % k-1 closest points
                    here
42                 currentPoint =
                    imageRepresentation(
                    rowInRepresentationMatrix, :)
                    ;
43

```

```

44         u = zeros(1, 5);
45         w = 0;
46
47         for t = 1:size(points, 2)
48             point =
49                 imageRepresentation(
50                     points(1, t), :);
51             w1 = exp(sum(((
52                 currentPoint - point)
53                 .^2) ./ sigmas))/
54                 sigmaProds;
55
56             u = u + w1*point;
57             w = w + w1;
58         end
59
60         u = u./w;
61
62         newImageRepresentation
63         (
64             rowInRepresentationMatrix
65             , :) = [
66             currentPoint
67             (1, 1),
68             currentPoint
69             (1, 2), u(1,
70             3), u(1, 4),
71             u(1, 5)];
72
73         end
74
75         imageRepresentation = newImageRepresentation;
76
77         end
78
79         % row = zeros(numPixels, 1);
80         % col = zeros(numPixels, 1);
81         % for i = 1:size(imageRepresentation, 1)
82         %     row(i, 1) = imageRepresentation(i, 1);
83         %     col(i, 1) = imageRepresentation(i, 2);
84         % end
85         %
86         % figure;
87         % plot(row, col, 'r*');
88         % hold on;
89
90         outputImage = zeros(size(processedImage));
91         for x = 1:size(processedImage, 1)
92             for y = 1:size(processedImage, 2)

```

```

78         rowInRepresentationMatrix = (x-1) * size
          (processedImage, 2) + y;
79     for i = 1:3
80         outputImage(x,y,i) =
          imageRepresentation(
          rowInRepresentationMatrix, i
          +2);
81     end
82     end
83     end
84
85 end

```

Implementation Details

We took a $K=200$ nearest neighbours and took their weighted average to determine the mean shift. Note that the weights were determined by a 5 dimensional Gaussian vector in space (X,Y) and colour (R,G,B) . This was repeated over N iterations for the image to converge.

Result Images

Original scaled



Mean shift segmented



Optimum Parameters

hs = 20

hr = 20 Number of Iterations = 20